

SPHERE App — User Handbook

What is SPHERE?

SPHERE (Synthetic Privacy-preserving Honest Evaluation and Release Engine) is the end-to-end platform for generating, certifying, and sharing privacy-preserving synthetic data:

1. **Generate** — produce a synthetic twin of your sensitive dataset
2. **Evaluate & Certify** — formally score fidelity and privacy, issue a certificate
3. **Share** — upload the synthetic data to Dropbox or Zenodo
4. **Post to Catalog** — list the dataset in the SPHERE World public directory
5. **Analyse with AI** — use Claude to develop and run analyses on the synthetic data, then apply the resulting code to your real dataset locally

One rule governs all five steps: **your real data never leaves your machine.**

Other things do leave, and you should know which. You choose when to upload a synthetic file (step 3) and when to post catalog metadata (step 4). Whenever you use the AI features (step 5), your conversation is routed through SPHERE's own servers and stored there — whether you are spending SPHERE credits or using your own Anthropic key.

Generating, evaluating and saving a certificate also need an internet connection. The work itself runs on your Mac, but before each of these steps starts, the app asks SPHERE's server to approve it. The request is tied to your account and says which step it is, the row count, column count and file size where the app knows them before the step starts, and the app version. If SPHERE can't be reached, the app keeps trying for up to about 30 seconds, then stops the step before any work is done and tells you why. And after each generate, evaluate, certify or share, the app sends SPHERE a small usage record, tied to your account: which step (and, for a share, where to), when it happened, how long it took (for generate and evaluate), row and column counts, file size and the app version. For generate and evaluate, the counts and file size are your original file's (a twin has the same number of rows and columns as the original); for a share they describe the synthetic file. Neither the approval request nor the record ever contains your data, file names, paths, column names or values. All of this is documented in *What SPHERE's servers receive* below.

This describes the account edition, which you sign in to. The licence-key (pilot) edition has no SPHERE account: it asks for no approvals, sends no usage records, and generates, evaluates and certifies offline as before.

Privacy Architecture

Understanding the privacy boundaries is central to using SPHERE confidently. This section covers all five panels at once so you know exactly what stays local and what leaves.

What Stays on Your Computer

Item	Why it stays local
Your original (real) data	Read only by SPHERE's own code on this Mac — to synthesise the twin, hash it and score it. Never uploaded, never transmitted, never sent to the model.
Synthesis algorithm	Runs entirely in the app process on your machine.
Evaluation computation	All statistical comparisons (fidelity, privacy tests) run locally.
SPHERE certificate	Generated locally, saved as an HTML file on your disk.
Dropbox and Zenodo tokens	Stored encrypted via macOS Keychain. Used only to talk to those services directly. Never sent to SPHERE's servers.
Your Anthropic API key	Stored encrypted via macOS Keychain. It is forwarded through SPHERE's server on each AI request so usage can be metered — it passes through and is not stored there.
Column names, dimensions & scores used in AI auto-fill	Derived from the synthetic file only — not the real one.
analysis.py and all session figures	Written and executed entirely inside an isolated local sandbox.

Generation, evaluation and certification run on your Mac, but in the account edition each needs an internet connection, because SPHERE's server approves it before it starts. The only things any of them sends are that approval request and the usage record, both described in the next section — never your data, file names, paths, column names or values.

What SPHERE's Servers Receive

Your account. Your email, name, institution or organization, organization type and role or title, and your acceptance of the SPHERE use agreement — see *Your Account and the SPHERE Use Agreement*.

The AI features do not talk to Anthropic directly. Every request is routed through SPHERE's own proxy at `api.sphereworld.ai`, which is how usage is metered against your account. This applies whether you are spending SPHERE credits **or** using your own Anthropic key.

That server receives, and **retains**:

- **The full conversation** — every message you send and every reply the model returns, including tool results and any real-data results you approved sending for the report.
- **Usage metadata** — token counts, model, timestamps, and your account identity.

It does **not** receive your real data, your real file's path, or your Anthropic key in storable form (the key is forwarded upstream on each request and not kept).

Sign-in is required to use the app, so this applies to all AI use. If your work is governed by an agreement that forbids conversation content leaving your institution, use the Generate, Evaluate and Share tabs and not the AI features — those tabs send SPHERE no conversation content and no data, only the approval requests and usage records described next.

Step approvals. Before a generate, an evaluate or a certificate save starts, the app asks SPHERE's server to approve it, and the server records the request at that moment, tied to your SPHERE account. The request contains only:

- **A random request ID**, made fresh for each step, so a request the app has to retry is recorded once
- **The step** — generate, evaluate or certify
- **Row count, column count and file size, where the app knows them before the step starts.** In practice that is:
 - *generate*: the size of your **original (real) file**
 - *evaluate*: the size of your **original (real) file**
 - *certify*: none
- **App version**

The server adds the time and your account, and sends back an approval signed by SPHERE, which the app checks before the step goes ahead. The request never contains your data, file names, paths, column names or values, and the step itself — reading your file, building the twin, scoring it, writing the certificate — still runs entirely on your Mac. Because the approval is recorded when it is given, a step that fails or is cancelled after it was approved still shows as approved.

A certificate save is when you save a certificate, or SPHERE AI saves one:

- **Generate Certificate** on the Evaluate tab, and **Save as ZIP** in the Share section (the ZIP holds the certificate). Both are approved after you choose where to save and before anything is written, so cancelling the save dialog sends nothing.

- The certificate SPHERE AI saves next to the twin when you load your original data (or an example) there. Doing so generates the twin, evaluates it and saves its certificate: three steps, each approved as above. If the certificate's approval fails, the certificate is not saved and SPHERE AI tells you why.

A certificate uploaded with a Dropbox or Zenodo share needs no approval of its own; the upload already needs a connection. Neither does the copy of it the app keeps in its share history. Saving a SPHERE AI session report is not a certificate and needs no approval.

If SPHERE can't be reached, or its server is briefly unavailable, the app keeps trying for up to about 30 seconds, so a brief drop in the connection does not stop you. The step stops before any work is done if there is still no approval after that, or straight away if SPHERE turns the request down, its answer does not check out, or your sign-in has expired. The app then says which: that it needs an internet connection, that your SPHERE sign-in has expired and you need to sign in again, or that SPHERE couldn't approve the step right now and you should try again in a moment.

The licence-key (pilot) edition asks for no approvals: it has no SPHERE account, and its generate, evaluate and certify work offline as before.

Usage records. After each successful generate, evaluate, certify or share, the app saves a small record on this Mac and sends it to SPHERE, tied to your SPHERE account, the next time it is online with that account signed in. It is always on for signed-in users; there is no setting to turn it off. Each record contains only:

- **A random record ID**, so a record that arrives twice is stored once
- **The step** — generate, evaluate, certify or share — and, for a share, its destination (Dropbox, Zenodo or the SPHERE World catalog). Saving a ZIP is recorded as a certify step, because the ZIP holds the certificate.
- **The time the step completed**, and for generate and evaluate, how long it took
- **Row count, column count and file size**, where they apply. They do not all describe the same file:
 - *generate*: the row count, column count and size of your **original (real) file**. The rows and columns are counted in the twin it produced, which has the same number of rows and columns as the original.
 - *evaluate*: the row count, column count and size of your **original (real) file**
 - *share*: the synthetic file you shared — its size for Dropbox or Zenodo, its row and column counts for the catalog
 - *certify*: none
- **App version**
- **The approval ID** for generate, evaluate and certify, linking the record to the step approval that came before it (a share has none)

It never contains your data, file names, paths, column names, cell values or dataset titles. Your account is identified by your sign-in, not by anything in the record.

In the account edition, generate, evaluate and certify need a connection, for the approval described above. Signing in needs a connection too: the first time, and again once 7 days have passed since the app last confirmed your sign-in with SPHERE. It does that when it starts with a connection and when you sign in; leaving the app open does not by itself restart the 7 days. Within those 7 days an app without a connection stays signed in, but it cannot generate, evaluate or certify until it can reach SPHERE again. A record that could not be sent straight away — for example because the connection dropped after the step was approved — waits on your Mac and is sent the next time the app is running with a connection and that account signed in: after your next step, at the next launch, or within 15 minutes. Sending the record happens in the background and never delays, blocks or fails the step itself; only the approval comes before the step.

Each time it sends, the app picks out only the records made by the account that is signed in at that moment. Signing out does not delete records that are still waiting: they stay in the app's data folder on this Mac and are sent if that account signs in again. A record still unsent after about three months is discarded, and if more than 10,000 are waiting, the oldest are discarded first.

What Is Uploaded (Synthetic Data Only)

When you share a dataset, **only the synthetic CSV** is uploaded — never the real data.

- **Dropbox:** Synthetic CSV and certificate go to your Dropbox account via OAuth. SPHERE has no access beyond the folder it creates.
- **Zenodo:** Synthetic CSV and certificate go to your Zenodo account via your personal API token.

What the SPHERE World Catalog Receives

When posting to the public catalog, the API receives metadata only — no data rows, ever:

- Dataset title, description, domain, sharing terms
- Public Dropbox/Zenodo link to the *synthetic* file
- Fidelity and privacy scores (aggregate numbers only)
- Whether the dataset was SPHERE-generated or externally produced
- Whether a certificate exists · Storage method (Dropbox or Zenodo)

What Claude Receives — Auto-Fill (Catalog Post)

The **Generate with AI** button in the Post panel sends a prompt to Claude (Anthropic API) containing:

1. The synthetic CSV **filename** (not its contents)
2. Column **names** from the synthetic header (up to 10, then "etc.")
3. **Dimensions** — row and column counts of the synthetic file
4. **Aggregate scores** — fidelity and privacy composites
5. Whether the dataset was SPHERE-generated or external

What Claude never receives: individual data records, cell values, your real data, or any information that could identify research participants. Note that column *names* are sent (item 2 above); in some domains a column name is itself sensitive.

What Claude Receives — SPHERE AI Panel

In the SPHERE AI tab, Claude operates an agent loop over a Python sandbox. It receives:

- The **contents of the synthetic CSV** (Claude reads it via Python tools, e.g. `pd.read_csv`)
- Conversation messages and tool results (stdout/stderr from local Python execution)
- Figures it requests to view
- Files you explicitly attach in the chat

What Claude receives, and what it does not

Claude never receives your real dataset. The file is never uploaded, and neither is its path. Every tool the agent can call — `read_file`, `list_files`, `write_file`, `view_figure` — resolves only inside that session's own directory, and during a deploy the real CSV is staged into a *separate* directory that the sandbox refuses to open or even list.

Individual records are not sent. The analysis is instructed, as a hard rule, to print aggregates only — counts, means, model coefficients, group sizes — and to describe data-quality problems by **column name, row position and count** rather than by showing the rows. If it needs to show what a problem row looks like, it uses the SPHERE twin, which is what the twin is for.

What is sent, and when:

Channel	When	What
Error traceback	Automatically, if <code>analysis.py</code> fails on real data — up to 3 attempts	File name, line number and exception class only. The exception message is dropped entirely , and source lines are never echoed.
Analysis output + figures	Only when you press Send this and write the report	The exact text shown to you on screen, plus the figures as rendered images.

Nothing else from a real run reaches Claude.

Before anything is sent, you see it

Pressing **Review what will be sent** does not send anything. It shows you the exact characters that would leave the machine, plus thumbnails of every figure. Only then does **Send this and write the report** transmit — and it transmits precisely the string you were shown, not a recomputed one.

Two filters run before that text is displayed, and both are described honestly:

- **Identifier scrub.** SPHERE reads your real CSV in its own process, collects values from identifier-like columns (MRNs, names, accession numbers), and removes any line containing one. Those values never enter the interface process. If your file has no such column — which is normal for de-identified extracts — the app tells you so and removes nothing.
- **Size cap.** Output is capped, and if the cap is reached the notice appears *inside* the text you review, so you can see that something was cut.

Sandbox enforcement

Code the agent writes never runs inside SPHERE. It runs as a separate frozen binary launched through `/usr/bin/sandbox-exec` in a macOS Seatbelt profile it cannot renegotiate. There is no unconfined fallback: if the bundled runtime is missing, the tool call fails rather than running unprotected. Inside that profile it:

- has **no network** — TCP, UDP, DNS and unix sockets are all denied
- has **no credentials** in its environment (an allow-list of eight variables)
- can **write** nowhere but its own session directory
- can **read** exactly one data file, named by literal path — the twin, or during a deploy the staged real copy
- cannot read your stored API key, other sessions' outputs, your home directory, cloud-synced folders, other applications' data, or external volumes

Child processes inherit the sandbox, and symlink, hard-link, firmlink and nested-sandbox escapes were each attempted and failed. `npm run test:boundary` runs these as a test suite against the real binary rather than asserting them.

Limits you should know about

Honest disclosure of what this design does **not** guarantee:

1. **Figures are sent as rendered images.** A per-subject plot shows one point per person. This is deliberate — it is how the report compares twin against real — but it means figures carry

individual-level structure. Review the thumbnails.

2. **The no-records rule is an instruction, not a lock.** The agent writes the analysis, so nothing inspects the script before it runs. The identifier scrub is a backstop and cannot recognise a table of per-subject numbers that prints no identifier. The review step is what actually enforces this — please read it.
 3. **The sandbox denies by name, not by allow-list.** Python needs parts of `~/Library` to start, so the profile permits by default and denies specific areas. Locations outside both the deny list and your data are readable.
 4. **SPHERE's own code does read your real file** — to stage it, hash it, synthesise the twin, and score privacy and fidelity. That is fixed application code the model never influences, and it is not sandboxed. It is, however, launched with a stripped environment: your API key and every other credential in the app's environment are withheld from it, so a bug in that code cannot turn into a network call on your account. "Claude never sees your real data" and "SPHERE never touches your real data" are different statements; only the first is true.
-

Your Account and the SPHERE Use Agreement

SPHERE needs a free account, whether you create one with an email address and password or sign in with Google or Microsoft. The desktop app and the website at app.sphereworld.ai use the same account.

Your profile. Every account has a name, an institution or organization, an organization type (academic, non-profit, government, industry or other) and a role or title. The sign-up form asks for them; a Google or Microsoft sign-in does not, so SPHERE asks for them once after you sign in. They are kept with your account, and SPHERE's administrators use them to see who uses SPHERE — for example, by organization type.

The use agreement. Before you use SPHERE you accept its use agreement once:

RESEARCH USE ONLY — NO COMMERCIAL USE. SPHERE and the data it produces are provided solely for non-commercial academic research. Commercial use is strictly prohibited, including training, fine-tuning, or developing any commercial machine-learning, artificial-intelligence, foundation, or generative model, and any for-profit product, service, or algorithm. Commercial use requires a separate written license.

Accepting it in the app or on the website counts for both. SPHERE records that your account accepted it, which version, when, and the profile details you confirmed. If the agreement changes, you are asked again.

In the account edition of the app (from 0.3.7), generate, evaluate and certificate saves are not approved until the agreement has been accepted. The pilot (licence-key) edition has no account and is not affected.

The SPHERE Workflow

All five tabs form a single linear workflow. You do not have to complete every step — use as many as are relevant to your work.

Step 1 — Generate Synthetic Data

Tab: Generate

1. Drag your real CSV into the app, or click to browse. The file is read locally — nothing is uploaded.
2. Configure generation parameters (synthetic rows, noise level, random seed).
3. Click **Generate**. In the account edition the app first asks SPHERE's server to approve the step, so you need an internet connection (see *Step approvals*); the pilot edition asks for no approval. The SPHERE algorithm then runs on your machine and produces a synthetic CSV.
4. Download the synthetic CSV to a location of your choice.

Your real data is only ever read from disk — it is never stored by the app or transmitted anywhere. In the account edition, the approval request carries only a random request ID, the step, your original file's size and the app version.

Step 2 — Evaluate & Certify

Tab: Evaluate, Certify & Share

1. Load your real CSV and the synthetic CSV to evaluate (SPHERE-generated or from any other tool).
2. Click **Evaluate**. In the account edition SPHERE's server approves the step first, so this needs an internet connection. The app then computes: - **Fidelity scores** — how closely the synthetic data matches the real data's statistics (mean, variance, correlations, KS distance). - **Privacy scores** — how resistant the synthetic data is to singling-out, linkability, and inference attacks.
3. Review the scores and the visual report.
4. Click **Generate Certificate** to create a signed HTML evaluation report, saved to your disk. In the account edition, SPHERE's server approves the save after you choose where to save it and before the file is written, so this also needs a connection.

All computation happens on your machine. No data leaves during this step: in the account edition, the approval requests carry only a random request ID, the step, your original file's size (for the evaluation) and the app version.

Step 3 — Share to Cloud

Tab: Evaluate, Certify & Share → Share section

Three options are available:

Option A: Upload to Dropbox Connect via OAuth (direct with Dropbox, not SPHERE). Click **Share to Dropbox** — the synthetic CSV and certificate are uploaded to a new folder and a public link is generated.

Option B: Upload to Zenodo Enter your Zenodo API token (stored locally via Keychain). Click **Upload to Zenodo** — the files are uploaded and a permanent DOI is generated.

Option C: Save as ZIP Save a ZIP of the synthetic data and certificate for manual distribution. Nothing is uploaded. Because the ZIP holds the certificate, in the account edition it counts as a certificate save: SPHERE's server approves it after you choose where to save it and before anything is written, so this needs a connection.

Step 4 — Post to SPHERE World Catalog

Tab: SPHERE World

Once your synthetic data is hosted on Dropbox or Zenodo, you can make it discoverable in the public catalog at sphere-world.vercel.app.

1. On your dataset card, click  **Post to SPHERE World**.
2. Fill in the title and description manually, or click **Generate with AI** to auto-fill. Set the domain and sharing terms: -  **Open (with citation)** — anyone can download; just cite your work. -  **Request access** — interested parties must email you first. -  **Revenue share** — commercial licensing terms apply.
3. Click **Post**. Only metadata and the public link are sent — no data rows.
4. The button changes to  **Update on SPHERE World**. Click it any time to change the title, description, or sharing terms. The same catalog record is updated in place — no duplicate is created.

Deduplication: Each dataset has a stable local UUID. Re-posting always updates the same catalog record. If you try to post a dataset whose title already exists in the catalog (posted by you or

someone else), you will be asked to choose a different title.

Step 5 — Analyse with SPHERE AI

Tab: SPHERE AI

SPHERE AI lets you build a rigorous data analysis by chatting with Claude — using the synthetic data as a safe stand-in throughout development. When the analysis is ready, the app runs it locally against the real data. **Claude is not involved in that final execution step at all.**

Setup

- **Anthropic API key** — obtained at `console.anthropic.com/settings/keys` . Entered once, stored encrypted via macOS Keychain.
- **Model** — Claude Opus 5 (most capable), Fable 5 (extended thinking), Sonnet 5 (balanced), Haiku 4.5 (fastest). Changed in the Settings panel.

How the session works

1. **Load a synthetic CSV.** Click the green *Synthetic* chip and pick your CSV, or click a built-in example. An isolated sandbox folder is created and the synthetic file is placed inside it.
2. **Load the real CSV (optional, needed for deploy).** Click the red *Real* chip. This path is stored in the app process only — it is never shared with Claude.
3. **Chat with Claude.** Type a prompt and press Enter (⇧Enter for new line). Built-in starter prompts: - "Summarize the dataset: column types, missing values, and key distributions." - "Identify the strongest associations between variables and quantify them with effect sizes and p-values." - "Propose three hypothesis-driven analyses, then run the most informative one."
4. **Claude writes and runs code.** The agent iteratively writes Python in `analysis.py` , installs packages (`matplotlib` , `seaborn` , `statsmodels` ...), and executes in the sandbox. Figures appear in the right panel. All execution is local.
5. **Iterate.** Ask follow-ups, change chart styles, add statistical tests. Claude refines `analysis.py` across turns. Click **Review analysis.py** to inspect the current script.
6. **Run on real data.** Once the synthetic run shows a green exit code, click **Deploy on real**. The app (not Claude): - Copies your real CSV into the sandbox as a temporary file. - Runs `analysis.py` locally as a subprocess against that copy. - Immediately deletes the copy when the script finishes. - Shows the real-data figures alongside the synthetic figures.

No message is sent to Claude during this step. Claude already finished its work on the synthetic data.

7. **Save the session report.** After a successful run, a self-contained HTML report is generated containing the full conversation transcript, `analysis.py`, both figure sets, and an AI-written narrative summary. Click **Open report in sandbox** to reveal it in Finder.

File attachments

Click **+** in the chat input bar to attach files (PDFs, images, text/code). Attached files are sent to the Claude API as conversation context. Do not attach files containing real personal data.

Tools Claude can use

Tool	What it does
<code>python</code>	Executes Python code in the sandbox. Output and errors are returned to Claude.
<code>pip_install</code>	Installs a Python package into the sandbox environment.
<code>read_file</code>	Reads a file from the sandbox directory.
<code>write_file</code>	Writes or overwrites a file in the sandbox directory.
<code>list_files</code>	Lists files present in the sandbox directory.
<code>view_figure</code>	Renders a saved figure so Claude can inspect and refine it.

Buttons reference

Button	What it does
Send	Sends your message to Claude. Press Enter (⇧Enter for new line).
 Stop	Interrupts Claude at the next turn boundary.
Review analysis.py	Shows the current state of the analysis script in a dialog.
Deploy on real	App runs <code>analysis.py</code> locally against your real CSV, then sends the <i>results</i> to Claude for the comparison section.
Sandbox folder	Opens the sandbox directory in Finder (scripts, figures, report).
 Reset	Clears the conversation and starts a fresh session with the same synthetic CSV.

First-launch pre-loading

The first time the SPHERE AI tab is used, `matplotlib`, `seaborn`, and `statsmodels` are pre-installed in the sandbox Python environment. A progress banner is shown (typically 30–90 seconds). Subsequent sessions start immediately.

Managing Your Datasets

Syncing and Deleting

Delete from Dropbox / Zenodo: Click *Delete from Dropbox* (or Zenodo) on any dataset card. The app deletes the remote folder, removes the local history entry, and automatically removes the catalog listing.

Sync: Click **Sync** in the SPHERE World tab to reconcile your local history with remote storage. If a file was deleted outside the app (e.g. directly in Dropbox's web UI), the sync detects it and removes the stale entry and catalog listing.

Updating a Catalog Listing

Click  **Update on SPHERE World** on any dataset card to change the title, description, domain, or sharing terms. The existing catalog record is updated — no duplicate is created.

SPHERE vs. External Synthetic

The app and catalog distinguish two dataset origins:

- **SPHERE** badge (cardinal red): The synthetic data was generated by the SPHERE algorithm built into this app.
- **External synthetic** badge (grey): The data was generated by a third-party tool (e.g., SDV, synthpop, CTGAN) and then evaluated using SPHERE's metrics.

This distinction is shown on dataset cards in both the app and the catalog.

Data Flow Summary

Your real data (CSV)

|

▼

[SPHERE App – on your Mac]

|

|→ Synthesis algorithm —————→ Synthetic CSV (saved locally)

|

|→ Evaluation & scoring (local)

|

|→ Certificate HTML (local)

|

|→ AI auto-fill (catalog post)

|

→ Claude API receives: filename,
column names, dimensions, scores

|

|→ SPHERE AI panel

|

|→ Synthetic CSV → sandbox

|

Claude writes analysis.py via python / pip / write_file tools

|

All execution: local subprocess on this Mac

|

Claude API receives: conversation + synthetic CSV contents only

|

|→ "Deploy on real" – app-only step, NO Claude API call

|

App copies real CSV → sandbox → runs analysis.py → deletes copy

|

Real data never transmitted; Claude is not involved in this step

|

|→ Approval before each generate / evaluate / certify (account edition)

|

→ sent to SPHERE before the step starts; no connection = no step

|

(tied to your account): request ID, step, app version, and

|

for generate/evaluate your real file's size

|

– no data, file names, paths, column names or values

|

|→ Usage record after each generate / evaluate / certify / share

|

→ saved on this Mac, sent to SPHERE when online and signed in

|

(tied to your account): record ID, step, destination, time,

|

duration, rows, columns, file size, app version, approval ID

|

(for generate/evaluate, rows, columns and size are your real file's)

|

– no data, file names, paths, column names or values

Frequently Asked Questions

Q: Why does SPHERE ask for my institution and role, and for me to accept an agreement?

SPHERE is free for non-commercial academic research. The agreement states that limit, and your profile lets SPHERE see who uses it. You provide them once: the app and the website share the same account, so accepting on one covers the other.

Q: Does SPHERE ever see my real data? Your real data is read on this Mac by SPHERE's own code — to build the twin, hash it and score it — and is never transmitted to SPHERE's servers, Anthropic, Dropbox, Zenodo or any other service. Three other things do reach SPHERE's servers, and none contains your data: your *conversation* with the AI, which is routed through SPHERE and retained there; and, in the account edition, the approval request sent before each generate, evaluate or certificate save and the small usage record sent after each generate, evaluate, certify or share. See *What SPHERE's servers receive*.

Q: Does SPHERE record how I use the app? Can I turn that off? Yes, in the account edition.

Before each generate, evaluate or certificate save, SPHERE's server approves the step and records the approval request — a random request ID, the step, the row count, column count and file size where the app knows them before the step (in practice your original file's size for generate and evaluate, nothing for a certificate) and the app version — tied to your account. After each successful generate, evaluate, certify or share, the app also saves a usage record on this Mac — a random record ID, the step, the share destination, when it happened, how long it took (generate and evaluate only), row and column counts, file size, app version and, except for a share, the ID of the step's approval — and sends it to SPHERE, tied to your account, the next time it is online with that account signed in. For generate and evaluate, the row count, column count and file size are your original file's (a twin has the same number of rows and columns as the original); for a share they describe the synthetic file. It is always on for signed-in users; there is no opt-out. Neither ever contains your data, file names, paths, column names, cell values or dataset titles.

Q: Do generate, evaluate and certify work without an internet connection? Not in the account edition. The work itself still runs on your Mac, but SPHERE's server has to approve each generate, evaluate and certificate save (including Save as ZIP) before it starts (see *Step approvals*), so each needs a connection. The app keeps trying for up to about 30 seconds to ride out a brief drop; if it still can't reach SPHERE, the step stops before any work is done and the app says it needs an internet connection. Uploading to Dropbox or Zenodo, posting to the catalog and the AI features have always needed a connection.

Signing in works as before: it needs a connection the first time, and again once 7 days have passed since the app last confirmed your sign-in with SPHERE (it does that when it starts with a connection and when you sign in; leaving the app open does not by itself restart the 7 days). If SPHERE starts while you are offline, you stay signed in, and the account menu says that generating, evaluating and

certifying need a connection and how many days are left before you must sign in again; a banner warns you in the last two days. Once they run out, the app opens at the ordinary sign-in screen the next time it starts, and you need a connection to get past it. A usage record that could not be sent straight away waits on your Mac until you are online and signed in to the same account.

The licence-key (pilot) edition is different: it has no SPHERE account, asks for no approvals, and generates, evaluates and certifies offline as before.

Q: Can the synthetic data be traced back to real individuals? SPHERE evaluates exactly this risk using three attack models: singling-out (can you identify a unique individual?), linkability (can you link records across datasets?), and inference (can you predict a sensitive attribute?). The privacy scores reflect how resistant the synthetic data is to each attack. Higher scores mean greater protection.

Q: Does Claude ever access my real data? No. In the SPHERE AI tab, Claude only ever sees the synthetic CSV and the code it writes itself. When you click "Deploy on real", the app — not Claude — copies your real CSV into the sandbox, runs `analysis.py` as a local subprocess, and immediately deletes the copy. No message is sent during the deploy itself. If you then press the report button, the aggregate results you have reviewed on screen are sent — that is a separate, explicit step.

Q: Is it safe to use SPHERE AI with sensitive datasets? While developing the analysis, Claude works only from the SPHERE twin, which by design contains no real individuals' records — so aim for good privacy scores before you rely on it. On a "Deploy on real" the script runs locally against your real file and Claude receives only the aggregate results, and only after you have read them on screen and pressed send. Your real dataset is never uploaded. See *Limits you should know about* above for what this does not guarantee.

Q: Can Claude access files outside the sandbox? No. Code the agent writes runs in a macOS Seatbelt sandbox with no network and no credentials, able to read exactly one data file — the twin, or during a deploy the staged real copy — and able to write only inside its own session directory. It cannot reach your home directory, cloud-synced folders (iCloud, OneDrive, Box, Dropbox), other applications' data, other SPHERE sessions, your stored API key, or external volumes. Child processes inherit the confinement. One honest caveat: the profile denies by name rather than permitting by allow-list, because Python needs parts of `~/Library` to start — so a location that is neither your data nor on the deny list is readable. `npm run test:boundary` attempts each of these escapes against the real binary.

Q: What exactly is sent to Anthropic's API? It does not go to Anthropic directly — it is routed through SPHERE's proxy, which meters and retains it (see *What SPHERE's servers receive*). The content is: conversation messages, tool results (stdout/stderr from local Python runs), figures Claude requests to view, and files you explicitly attach. The synthetic CSV's contents can be read by

Claude via `pd.read_csv` inside the `python` tool. Do not use synthetic data with residual real values if you have concerns about Claude reading column contents.

In the catalog post panel (AI auto-fill): only the filename, column names, dimensions, and aggregate scores from the synthetic file.

Q: Who can access the data on Dropbox or Zenodo? That depends on the sharing terms you set in the catalog. "Open (with citation)" means anyone with the link can access the synthetic file. "Request access" means interested parties must email you first.

Q: What happens if I delete the dataset from Dropbox but it's in the catalog? The next time you run **Sync** in the app, the stale entry is detected and the catalog listing is removed automatically.

Q: Can I update my catalog listing after posting? Yes. Click  **Update on SPHERE World** on your dataset card at any time. The existing catalog entry is updated in place — no duplicate is created.

Q: Are my API keys and tokens stored securely? All keys are stored in your macOS user data directory, encrypted via macOS Keychain where supported. Your Dropbox and Zenodo tokens are never sent to SPHERE's servers. Your Anthropic key is different: it is forwarded through SPHERE's proxy on each AI request so usage can be metered. It passes through and is not stored there.

Q: Can I evaluate a synthetic dataset generated by another tool? Yes. Use the Evaluate tab with any real CSV and any synthetic CSV, regardless of how the synthetic data was produced. It will be labelled "External synthetic" in the catalog.

Quick Reference

Action	Where	Privacy impact
Generate synthetic data	Generate tab	Runs locally. Account edition: needs a connection — SPHERE approves the step first (request ID, step, file size, app version; no data) — and a usage record (no data) is sent to SPHERE afterwards
Evaluate fidelity / privacy	Evaluate tab	Runs locally. Account edition: needs a connection — SPHERE approves the step first (request ID, step, file size, app version; no data) — and a usage record (no data) is sent to SPHERE afterwards

Action	Where	Privacy impact
Generate certificate	Evaluate tab	Written locally. Account edition: needs a connection — SPHERE approves the save first (request ID, step, app version; no data) — and a usage record (no data) is sent to SPHERE afterwards
Upload to Dropbox / Zenodo	Share section	Synthetic data and certificate only; account edition: usage record (no data) sent to SPHERE
Save as ZIP	Share section	Written locally; nothing uploaded. The ZIP holds the certificate, so in the account edition it needs a connection — SPHERE approves it first, like a certificate save (request ID, step, app version; no data) — and a usage record (no data) is sent to SPHERE afterwards
Post to SPHERE World catalog	SPHERE World tab → Post	Metadata + public link only; account edition: usage record sent to SPHERE
AI auto-fill description	Post panel → Generate with AI	Filename, column names, dimensions, scores sent to Claude
Update catalog listing	SPHERE World tab → Update	Same as Post — metadata only
Delete dataset	Card → Delete button	Removes from cloud storage + catalog
Sync	SPHERE World tab → Sync	Checks remote storage, removes stale entries
SPHERE AI — chat session	SPHERE AI tab	Synthetic CSV + conversation, via SPHERE's proxy — conversation retained
SPHERE AI — Deploy on real	SPHERE AI → Deploy on real	Analysis runs locally; results are sent only if you press the report button and approve them

Action	Where	Privacy impact
Session report	SPHERE AI → Open report	100% local — HTML saved to sandbox folder